

1 Description

We have now added SER¹ transactions support on top of Quelea [2]. The added mechanism uses an external key-value database etcd² which implements the Raft consensus protocol to offer a high performance distributed (but strongly consistent) store. I initially tried to use Quelea’s own SC operations to implement a simple locking mechanism. SC operations in Quelea are based on Cassandra’s CAS (compare and swap) lightweight transactions. However, implementation of CAS in Cassandra is very bad and show incorrect behavior even under light pressure (and totally breaks down under congestion).

In order to show the performance gain from assigning fine-tuned isolation guarantees to traditional SQL transactions, I implemented a very simple version of TPC-C in Quelea. For this experiment we did not need a full-fledged application with program logic details to show the performance gain from weakly isolated transactions. I created a set of simple transactions, according to TPC-C specification, that work with Quelea’s LWW registers. The transactions, roughly follow the defined TPC-C behavior. For example, our newOrder updates 6 registers each corresponding one task from the original TPC-C. This has to be extended to a full programming framework with an interesting enough logic, but for now, it suffices to support our claims. The transactions are executed according to the following table³:

new order	SER	5%
delivery	SER	40%
order status	MAV	5%
payment	MAV	45%
stock level	MAV	5%

2 Results

The following results are from a 3-node amazon ec2 cluster running (Cassandra/Quelea + etcd). The cluster consists of t2.medium instances at Oregon, Ohio, and Virginia. Note that the extremely high latencies are due to the

¹Traditional fully isolated transactions

²<https://coreos.com/etcd/>

³The isolation levels are assigned following the results in [1]

fact that all transactions are being coordinated since we are assuming a single object for now. This can be relaxed when we implement an object based lock and only coordinate transactions accessing the same object (e.g. in [2] experiments are on objects that are randomly chosen from a set of 10000, which rarely causes any blocking). The results show almost 50% lower latency when using SER transactions only when they are absolutely needed:




Figure 1: Latency vs number of concurrent clients

3 To Do

1. In the current lock implementation, clients try acquiring the lock and if they fail, they try again after 0.05 seconds. Although this cannot affect the average latency of all operations (which we are measuring), but it can potentially cause starvation for some clients and should be extended to a fair locking system later.

2. Currently, SER is implemented at the client side by acquiring a global lock. This task should be pushed into the server side in a more realistic system. We should also switch to an object-based lock to allow a higher through put.
3. Ultimately, we need a general purpose and full-fledged programming framework for writing EC applications on Quelea. The framework should be equal to L3 from the previous document.

References

- [1] P. Bailis, A. Davidson, A. Fekete, A. Ghodsi, J. M. Hellerstein, and I. Stoica. Highly available transactions: Virtues and limitations. *Proc. VLDB Endow.*, 7(3):181–192, Nov. 2013.
- [2] K. Sivaramakrishnan, G. Kaki, and S. Jagannathan. Declarative programming over eventually consistent data stores. In *Proceedings of the 36th ACM SIGPLAN Conference on Programming Language Design and Implementation*, PLDI '15, pages 413–424, New York, NY, USA, 2015. ACM.